

Rcpp and CUDA backend Logistic Regression

CS628 Parallism Algorithm

Xin Zhou

May 6, 2015

- 1 Intro
- 2 Logistic Regression Algorithm
- 3 Result

Outline

- 1 Intro
- 2 Logistic Regression Algorithm
- 3 Result

Background

The Rcpp and CUDA provide give us the possibility that we can wrap C++ code or CUDA code in R Package. Therefore, building a parallel framework for machine learning problem will improve the performance of several algorithm of certain R Package

Existed Parallel Framework

There are several packages combine CUDA or MPI into R Packages

- **Rmpi**

Existed Parallel Framework

There are several packages combine CUDA or MPI into R Packages

- **Rmpi**
- **R/parallel**

Existed Parallel Framework

There are several packages combine CUDA or MPI into R Packages

- **Rmpi**
- **R/parallel**
- **Multicore**

Existed Parallel Framework

There are several packages combine CUDA or MPI into R Packages

- **Rmpi**
- **R/parallel**
- **Multicore**
- **gputools**

Existed Parallel Framework

There are several packages combine CUDA or MPI into R Packages

- **Rmpi**
- **R/parallel**
- **Multicore**
- **gputools**
- **CudaBayesreg**

Existed Parallel Framework

There are several packages combine CUDA or MPI into R Packages

- **Rmpi**
- **R/parallel**
- **Multicore**
- **gputools**
- **CudaBayesreg**
- **RCUDA**

Rcpp Solution

with Rcpp we can integrate the .cu Code directly into package's dynamic linking library, and the special Makefile in Rcpp Package is a Makevars

```
PKG_LIBS = $(shell $(R_HOME)/bin/Rscript -e "Rcpp::LdLibraryPaths()"
CUDA_INCS = -I/usr/local/cuda/include -lcublas
```

```
GLMOBJECTS=\
./GLM/*.o \
...
```

```
RCPPGLM = RcppGLM.o
```

```
RCPP_EXPORT = RcppExports.o
```

```
OBJECTS = $(GLMOBJECTS) $(RCPPGLM) $(RCPP_EXPORT)
```

Outline

- 1 Intro
- 2 Logistic Regression Algorithm
- 3 Result

multinomial logistic regression problem

In the multinomial logistic regression problem, we have the probability definition of each class on each observation.

$$\ln\left(\frac{\pi_{ij}}{\pi_{iJ}}\right) = X\beta_j$$

multinomial logistic regression problem

- Here β_j means the j th column of matrix β .

multinomial logistic regression problem

- Here β_j means the j th column of matrix β .
- Since $\sum_{j=1}^J \pi_{ij} = 1$, which means that for each response Y_i , the summed probability in predictors space is 1.

multinomial logistic regression problem

- Here β_j means the j th column of matrix β .
- Since $\sum_{j=1}^J \pi_{ij} = 1$, which means that for each response Y_i , the summed probability in predictors space is 1.
- In Logistical regression, since we have N obervation \mathbf{Y} , and we use \mathbf{Y}_i represents each observation $(0, 0, 0, ..1, ..0)$.

multinomial logistic regression problem

- Here β_j means the j th column of matrix β .
- Since $\sum_{j=1}^J \pi_{ij} = 1$, which means that for each response Y_i , the summed probability in predictors space is 1.
- In Logistical regression, since we have N obervation \mathbf{Y} , and we use \mathbf{Y}_i represents each observation $(0, 0, 0, ..1, ..0)$.
- Therefore, we can define the probability of all classes:

probability of multinomial logistic regression problem

$$\pi_{ij} = \frac{e^{x_i \beta_j}}{1 + \sum_{j=1}^{J-1} e^{x_i \beta_j}} \quad j \neq J$$

$$\pi_{iJ} = \frac{1}{1 + \sum_{j=1}^{J-1} e^{x_i \beta_j}} \quad j = J$$

object function and gradient

$$L(\beta) = - \sum_{i=1}^N [\sum_{j=1}^{J-1} \mathbf{Y}_{ij} \sum_{k=1}^p X_{ik} \beta_{kj} - n_i [\log[1 + \sum_{j=1}^{J-1} e^{\sum_{k=1}^p X_{ik} \beta_{kj}}]]]$$

$$\nabla = - \sum_{i=1}^N [\mathbf{Y}_{ij} X_{ik} - n_i X_{ik} \pi_{ij}] = -X^T (Y - \pi)$$

In one vs all model

We also can define the Hessian matrix for the update:

$$\beta^{(t+1)} = \beta^{(t)} - \alpha H^{-1} \nabla.$$

And in multinomial model, for each class j , we will have

$$\mathbf{H}_{kk'} = \sum_{i=1}^N n_i X_{ik} X_{ik'} \pi_{ij} (1 - \pi_{ij})$$

Parallel

As the result, in gradient descent, we can paralyze them in several parts:

- Summed Loss function L : distribute the $[\log[1 + \sum_{j=1}^{J-1} e^{\sum_{k=1}^p x_{ik}\beta_{kj}}]]$ into each threads, and use **reduce to sum them**.

Parallel

As the result, in gradient descent, we can paralyze them in several parts:

- Summed Loss function L : distribute the $[\log[1 + \sum_{j=1}^{J-1} e^{\sum_{k=1}^p x_{ik}\beta_{kj}}]]$ into each threads, and use **reduce to sum them**.
- Matrix Multiplication: in gradient function ∇ , calculate the $-X^T(Y - \pi)$

Outline

- 1 Intro
- 2 Logistic Regression Algorithm
- 3 Result**

Classification result

2	1	6	1	5	4	9	8	5	7
5	2	6	3	3	4	7	2	5	7
8	3	3	4	9	0	6	4	0	5
5	3	3	4	9	0	6	4	0	5
9	1	2	4	3	9	6	3	0	4
8	1	2	4	3	9	6	4	0	5
4	8	3	9	8	2	8	2	1	6
4	3	3	4	8	2	8	2	1	6
2	0	7	2	1	7	2	3	0	1
2	9	7	4	1	7	7	7	0	1
6	1	6	6	1	7	2	2	2	7
6	1	6	6	1	7	2	2	2	7
9	8	8	9	4	8	0	3	1	6
9	8	8	4	4	5	6	3	1	6
7	2	7	5	1	8	8	8	3	8
7	2	7	5	1	8	8	8	3	8
1	7	5	5	7	0	8	8	3	7
1	7	5	5	7	0	8	8	3	7
0	2	3	1	5	6	4	0	4	9
0	2	3	1	5	6	4	0	4	9

Speed up on CUDA about 8x and Classification error ~ 0.13

